

# futex描述

futex()系统调用提供了一种方法用于等待某个特定条件的发生。一种典型的应用是在共享内存同步中作为阻塞装置。当使用futex时，绝大部分同步操作都在用户态完成，一个用户态程序只有在有可能阻塞一段较长的时间等待条件的发生时才使用futex。其他的futex操作可以用来唤醒等待特定条件发生的任何进程或线程。

一个futex是一个32位的值，它的地址通过futex()传入，futexs在所有的平台上都是32位的(包括64位系统上)，所有的futex操作都是在这个值上。为了在多个进程间共享futex，futex位于由mmap或者shmat创建的一块共享内存上(这种情况下，futex值在不同的进程上可能位于不同的虚拟地址，但这些虚拟地址都指向相同的物理地址)。在多线程程序中，将futex值放到一个所有线程共享的全局变量中就可以了。

当执行一个futex操作请求阻塞一个线程时，只有在 *uaddr == val* 时，内核才会执行阻塞操作，这个操作中的所有步骤(1. 导入uaddr的值; 2. 比较; 3.阻塞线程)将原子的执行，并且当其他线程在同一个futex值上并行操作时所有的步骤不会乱序。

futex的一种使用方式是用来实现锁，锁的状态(acquired or not acquired)可以用在共享内存中的原子标志表示。在非竞争的情况下，线程可以通过原子操作访问和修改锁的状态(这些操作全部在用户态操作，内核不会保存任何关于锁的状态)。从另一方面来说，另一个线程可能无法获取锁(因为锁已经被某个线程获取)，这个线程将通过如下这种方式执行futex()等待操作：

```
atomic<int> lock; // lock : 0. 锁未被获取 1.锁被获取
futex(&lock, FUTEX_WAIT, 1, NULL, NULL, NULL);
```

futex(FUTEX\_WAIT)将会检测lock的值，只有在等于1时才阻塞线程。当线程释放锁时，该线程必须首先重置锁的状态，然后执行futex操作唤醒阻塞在lock上的线程。

注意使用futex并没有显示的初始化和销毁操作，内核只有在一个指定的futex值上执行futex操作时(例如FUTEX\_WAIT)才会维护futex数据。

## 参数

uaddr指向futex值，在所有的平台上，futex值是一个4字节的整数并且必须4字节对齐。

对于某些阻塞操作，timeout参数是一个指向timespec结构的指针，表明了操作的超时时间。然而，在其他的某些操作下，它的最低4字节被作为一个整数值，这个整数值含义因futex操作的不同而不同，对于这些操作来说，内核会将timeout值转换为unsigned long，然后转换为uint32\_t，在接下来的说明中，它将表示为val2。

在需要的时候，uaddr2是指向第二个futex值的指针，val3的解释将依赖于具体的操作

## 操作

**FUTEX\_CLOCK\_REALTIME** (since Linux 2.6.28)

这个标志选项只能用在**FUTEX\_WAIT\_BITSET** 和 **FUTEX\_WAIT\_REQUEUE\_PI** 操作中，如果设置了这个标志选项，内核会将timeout视作基于**CLOCK\_REALTIME** 的绝对时间。如果没有设置这个标志，则内核将它视作基于**CLOCK\_MONOTONIC**的相对时间。

## FUTEX\_WAIT (since Linux 2.6.0)

本操作会检查uaddr是否等于val，如果相等的话，则睡眠等待在uaddr上的**FUTEX\_WAKE** 操作，如果线程开始睡眠，则它被认为是这个futex值上一个等待者。如果两者不相等，则操作失败并返回 *error EAGAIN*。比较uaddr和val的目的是为了避免丢失唤醒。

如果timeout参数不为NULL，则它指定了等待的超时时间(根据**CLOCK\_MONOTONIC**测量得出)。

## FUTEX\_WAKE (since Linux 2.6.0)

本操作会唤醒最多val个等待者，绝大部分情况下，val的值为1(只唤醒一个等待者)或INT\_MAX(唤醒所有的等待者)。**注意没有任何机制保证一定唤醒某些特定的等待者(比如被认为是高优先级的等待者)**

参数timeout, uaddr2, val3将被忽略。

# 代码

```
fn futex_wait(
    vaddr: VirtAddr,
    futex_op: i32,
    val: u32,
    timeout: usize,
) -> Result<usize, SyscallError> {
    let mut to = false;
    let current_task = current_task();
    let deadline = if timeout != 0 {
        Some(Duration::from_nanos(timeout as u64) + axhal::time::current_time())
    } else {
        None
    };
    loop {
        let key = get_futex_key(vaddr, futex_op);
        let process = current_process();
        if process.manual_alloc_for_lazy(vaddr).is_ok() {
            let real_futex_val = unsafe { (vaddr.as_usize() as *const
u32).read_volatile() };
            info!("real val: {:#x}, expected val: {:#x}", real_futex_val, val);
            if real_futex_val != val {
                return Err(SyscallError::EAGAIN);
            }
            futex_quque(key, &current_task, val);

            if let Some(deadline) = deadline {
                let now = axhal::time::current_time();
                to = deadline < now;
            }
            if timeout == 0 || !to {
                yield_now_task();
            }
            // If we were woken (and unqueued), we succeeded, whatever.
            // TODO: plist_del, not just iterate all the list
            if !futex_unqueue(key, &current_task) {
```

```

        return Ok(0);
    }
    if to {
        return Err(SyscallError::ETIMEDOUT);
    }
    // we expect signal_pending(current), but we might be the victim
    // of a spurious wakeup as well.
    #[cfg(feature = "signal")]
    if process.have_signals().is_some() {
        // 被信号打断
        return Err(SyscallError::EINTR);
    }
} else {
    return Err(SyscallError::EFAULT);
}
}
}

fn futex_wait_bitset(
    vaddr: VirtAddr,
    futex_op: i32,
    val: u32,
    timeout: usize,
    bitset: u32,
) -> Result<usize, SyscallError> {
    let mut to = false;
    let current_task = current_task();
    let deadline = if timeout != 0 {
        Some(Duration::from_nanos(timeout as u64) + axhal::time::current_time())
    } else {
        None
    };
    loop {
        let key = get_futex_key(vaddr, futex_op);
        let process = current_process();
        if process.manual_alloc_for_lazy(vaddr).is_ok() {
            let real_futex_val = unsafe { (vaddr.as_usize() as *const
u32).read_volatile() };
            info!("real val: {:#x}, expected val: {:#x}", real_futex_val, val);
            if real_futex_val != val {
                return Err(SyscallError::EAGAIN);
            }
            futex_quque(key, &current_task, val);

            // Check the bitset condition
            if (real_futex_val & bitset) == bitset {
                return Ok(0);
            }

            if let Some(deadline) = deadline {
                let now = axhal::time::current_time();
                to = deadline < now;
            }
            if timeout == 0 || !to {
                yield_now_task();
            }
        }
    }
}

```

```

    }
    // If we were woken (and unqueued), we succeeded, whatever.
    // TODO: plist_del, not just iterate all the list
    if !futex_unqueue(key, &current_task) {
        return Ok(0);
    }
    if to {
        return Err(SyscallError::ETIMEDOUT);
    }
    // we expect signal_pending(current), but we might be the victim
    // of a spurious wakeup as well.
    #[cfg(feature = "signal")]
    if process.have_signals().is_some() {
        // 被信号打断
        return Err(SyscallError::EINTR);
    }
} else {
    return Err(SyscallError::EFAULT);
}
}
}

fn futex_wake(vaddr: VirtAddr, futex_op: i32, val: u32) -> Result<usize, SyscallError> {
    let mut ret = 0;
    let key = get_futex_key(vaddr, futex_op);
    // 当前任务释放了锁，所以不需要再次释放
    let mut futex_wait_task = FUTEX_WAIT_TASK.lock();
    if futex_wait_task.contains_key(&key) {
        let wait_list = futex_wait_task.get_mut(&key).unwrap();
        // info!("now task: {}", wait_list.len());
        while let Some((task, _)) = wait_list.pop_front() {
            // 唤醒一个正在等待的任务
            info!("wake task: {}", task.id().as_u64());
            WAIT_FOR_FUTEX.notify_task(&task);
            ret += 1;
            if ret == val {
                break;
            }
        }
    }
    drop(futex_wait_task);
    yield_now_task();
    Ok(ret as usize)
}

fn futex_wake_bitset(
    vaddr: VirtAddr,
    futex_op: i32,
    val: u32,
    bitset: u32,
) -> Result<usize, SyscallError> {
    let mut ret = 0;
    let key = get_futex_key(vaddr, futex_op);
    let mut futex_wait_task = FUTEX_WAIT_TASK.lock();

```

```

        if let Some(wait_list) = futex_wait_task.get_mut(&key) {
            for (task, task_bitset) in wait_list.iter() {
                let wakeup = (val == 0 || ret < val) && ((*task_bitset & bitset) !=
0);

                if wakeup {
                    WAIT_FOR_FUTEX.notify_task(&task);
                    ret += 1;
                }
                if val != 0 && ret >= val {
                    break;
                }
            }

            wait_list.retain(|(_, task_bitset)| (*task_bitset & !bitset) != 0);
            if wait_list.is_empty() {
                futex_wait_task.remove(&key);
            }
        }
        drop(futex_wait_task);
        Ok(ret as usize)
    }

    /// To do the futex operation
    ///
    /// It may create, remove the futex wait task or requeue the futex wait task
    pub fn futex(
        vaddr: VirtAddr,
        futex_op: i32,
        val: u32,
        timeout: usize,
        vaddr2: VirtAddr,
        val2: usize,
        val3: u32,
    ) -> Result<usize, SyscallError> {
        let flag = FutexFlags::new(futex_op);
        match flag {
            FutexFlags::Wait => futex_wait(vaddr, futex_op, val, timeout),
            FutexFlags::Wake => futex_wake(vaddr, futex_op, val),
            FutexFlags::Requeue => {
                futex_requeue(val, val2, vaddr, vaddr2);
                Ok(0)
            }
            FutexFlags::WaitBitset | FutexFlags::RealTime => {
                futex_wait_bitset(vaddr, futex_op, val, timeout, val3)
            }
            FutexFlags::WakeBitset => futex_wake_bitset(vaddr, futex_op, val, val3),
            _ => Err(SyscallError::EINVAL),
        }
    }

    /// # Arguments
    /// * vaddr: usize
    /// * futex_op: i32
    /// * futex_val: u32
    /// * time_out_val: usize

```

```

/// * vaddr2: usize
/// * val3: u32
pub fn syscall_futex(args: [usize; 6]) -> SyscallResult {
    let vaddr = args[0];
    let futex_op = args[1] as i32;
    let futex_val = args[2] as u32;
    let time_out_val = args[3];
    let vaddr2 = args[4];
    let val3 = args[5] as u32;
    let process = current_process();
    let timeout = if time_out_val != 0 &&
process.manual_alloc_for_lazy(time_out_val.into()).is_ok()
    {
        let time_sepc: TimeSecs = unsafe { *(time_out_val as *const TimeSecs) };
        time_sepc.turn_to_nanos()
    } else {
        // usize::MAX
        0
    };
    // 释放锁，防止任务无法被调度
    match futex(
        vaddr.into(),
        futex_op,
        futex_val,
        timeout,
        vaddr2.into(),
        time_out_val,
        val3,
    ) {
        ok(ans) => Ok(ans as isize),
        Err(errno) => Err(errno),
    }
}

```

## 相关截图

---

C mlock.c

C futex\_wait\_timeout.c

C syscalls.c 3 X

C futetest.h 2

C futex.h ...\futex

C

linux-master &gt; kernel &gt; futex &gt; C syscalls.c &gt; do\_futex(u32 \_\_user \*, int, u32, ktime\_t \*, u32 \_\_user \*, u32, u32)

```
80
81     return ret;
82 }
83
84 long do_futex(u32 __user *uaddr, int op, u32 val, ktime_t *timeout,
85              u32 __user *uaddr2, u32 val2, u32 val3)
86 {
87     unsigned int flags = futex_to_flags(op);
88     int cmd = op & FUTEX_CMD_MASK;
89
90     if (flags & FLAGS_CLOCKRT) {
91         if (cmd != FUTEX_WAIT_BITSET &&
92             cmd != FUTEX_WAIT_REQUEUE_PI &&
93             cmd != FUTEX_LOCK_PI2)
94             return -ENOSYS;
95     }
96
97     switch (cmd) {
98     case FUTEX_WAIT:
99         val3 = FUTEX_BITSET_MATCH_ANY;
100         fallthrough;
101     case FUTEX_WAIT_BITSET:
102         return futex_wait(uaddr, flags, val, timeout, val3);
103     case FUTEX_WAKE:
104         val3 = FUTEX_BITSET_MATCH_ANY;
105         fallthrough;
106     case FUTEX_WAKE_BITSET:
107         return futex_wake(uaddr, flags, val, val3);
108     case FUTEX_REQUEUE:
109         return futex_requeue(uaddr, flags, uaddr2, flags, val, val2, NULL, 0);
110     case FUTEX_CMP_REQUEUE:
111         return futex_requeue(uaddr, flags, uaddr2, flags, val, val2, &val3, 0);
112     case FUTEX_WAKE_OP:
113         return futex_wake_op(uaddr, flags, uaddr2, val, val2, val3);
114     case FUTEX_LOCK_PI:
115         flags |= FLAGS_CLOCKRT;
116         fallthrough;
117     case FUTEX_LOCK_PI2:
118         return futex_lock_pi(uaddr, flags, timeout, 0);
119     case FUTEX_UNLOCK_PI:
120         return futex_unlock_pi(uaddr, flags);
121     case FUTEX_TRYLOCK_PI:
122         return futex_lock_pi(uaddr, flags, NULL, 1);
123     case FUTEX_WAIT_REQUEUE_PI:
124         val3 = FUTEX_BITSET_MATCH_ANY;
125         return futex_wait_requeue_pi(uaddr, flags, val, timeout, val3,
126                                     uaddr2);
127     case FUTEX_CMP_REQUEUE_PI:
128         return futex_requeue(uaddr, flags, uaddr2, flags, val, val2, &val3, 1);
129     }
130     return -ENOSYS;
131 }
```

```
apps > c > futex > C futex.c > main()
```

```
1  #include <unistd.h>
2  #include <sys/syscall.h>
3  #include <linux/futex.h>
4
5  int main() {
6      int futex_var = 0;
7
8      // 创建一个 futex, 初始值为 0
9      int *futex_ptr = &futex_var;
10
11     // 唤醒等待在 futex 上的线程（私有唤醒）
12     syscall(SYS_futex, futex_ptr, FUTEX_WAKE_PRIVATE, 1, NULL, NULL, 0);
13
14     return 0;
15 }
16
```

```
1  /*! Time structures.
2
3  The `time` module contains structures used to represent both
4  absolute and relative time.
5
6  - [Instant] is used to represent absolute time.
7  - [Duration] is used to represent relative time.
8
9  [Instant]: struct.Instant.html
10 [Duration]: struct.Duration.html
11 */
```

match返回类型必须相同，否则编译器会报错

枚举返回的结果类型，在后面的判断中比较大小时又出现类型不同的问题

```
let deadline: () = if timeout != 0 {
    enum MyResult {
        InstantResult(Instant),
        DurationResult(Duration),
        NoneResult,
    }

    let deadline: Option<MyResult> = if timeout != 0 {
        let result: Option<MyResult> = Some(match futex_op {
            FUTEX_CLOCK_REALTIME: i32 => MyResult::InstantResult(Instant::from_micros(timeout as i64) + axhal::time::current_time().into()),
            FUTEX_WAIT_BITSET: i32 => MyResult::DurationResult((axhal::time::current_time() + Duration::from_nanos(timeout as u64)).into()),
            _ => MyResult::NoneResult,
        });
        result
    } else {
        None
    };
};
```





```

fn futex_queue_with_bitset(key: FutexKey, curr: &CurrentTask, val: u32, bitset:
u32) {
    let mut futex_wait_bitset_task = FUTEX_WAIT_BITSET_TASK.lock();
    let wait_bitset_list = futex_wait_bitset_task.entry(key).or_default();
    wait_bitset_list.push_back((curr.as_task_ref().clone(), val, bitset));
}

```

```

fn futex_unqueue_with_bitset(key: FutexKey, curr: &CurrentTask, bitset: u32) ->
bool {
    let mut futex_wait_bitset_task = FUTEX_WAIT_BITSET_TASK.lock();
    if futex_wait_bitset_task.contains_key(&key) {
        let wait_bitset_list = futex_wait_bitset_task.get_mut(&key).unwrap();
        if let Some(index) = wait_bitset_list
            .iter()
            .position(|(task, _, b)| task.id() == curr.id() && *b == bitset)
        {
            wait_bitset_list.remove(index);
            return true;
        }
    }
    false
}

```

```

fn futex_wait_bitset(
    vaddr: VirtAddr,
    futex_op: i32,
    val: u32,
    timeout: usize,
    bitset: u32,
) -> Result<usize, SyscallError> {
    let mut to = false;
    let current_task = current_task();
    let deadline = if timeout != 0 {
        Some(Duration::from_nanos(timeout as u64) + axhal::time::current_time())
    } else {
        None
    };
    };

    loop {
        let key = get_futex_key(vaddr, futex_op);
        let process = current_process();

        if process.manual_alloc_for_lazy(vaddr).is_ok() {
            let real_futex_val = unsafe { (vaddr.as_usize() as *const
u32).read_volatile() };
            info!("real val: {:#x}, expected val: {:#x}", real_futex_val, val);

            if real_futex_val != val {
                return Err(SyscallError::EAGAIN);
            }
        }
    }
}

```

```

        futex_queue_with_bitset(key, &current_task, val, bitset);

    if let Some(deadline) = deadline {
        let now = axhal::time::current_time();
        to = deadline < now;
    }

    if timeout == 0 || !to {
        yield_now_task();
    }

    if !futex_unqueue_with_bitset(key, &current_task, bitset) {
        return Ok(0);
    }

    if to {
        return Err(SyscallError::ETIMEDOUT);
    }

    #[cfg(feature = "signal")]
    if process.have_signals().is_some() {
        return Err(SyscallError::EINTR);
    }
} else {
    return Err(SyscallError::EFAULT);
}
}
}

```

未添加futex\_wait\_bitset之前执行ZLM的结果为

```

1970-01-01 00:00:14.663 I [MediaServer] [9-MediaServer] Factory.cpp:41 registerPlugin | Load codec: H264
1970-01-01 00:00:14.786 I [MediaServer] [9-MediaServer] Factory.cpp:41 registerPlugin | Load codec: H265
1970-01-01 00:00:14.797 I [MediaServer] [9-MediaServer] Factory.cpp:41 registerPlugin | Load codec: JPEG
1970-01-01 00:00:14.808 I [MediaServer] [9-MediaServer] Factory.cpp:41 registerPlugin | Load codec: mpeg4-generic
1970-01-01 00:00:14.819 I [MediaServer] [9-MediaServer] Factory.cpp:41 registerPlugin | Load codec: opus
1970-01-01 00:00:14.830 I [MediaServer] [9-MediaServer] Factory.cpp:41 registerPlugin | Load codec: PCMA
1970-01-01 00:00:14.841 I [MediaServer] [9-MediaServer] Factory.cpp:41 registerPlugin | Load codec: PCMU
1970-01-01 00:00:14.843 I [MediaServer] [9-MediaServer] Factory.cpp:41 registerPlugin | Load codec: L16
9:
9:      transferring control: ./MediaServer
9:
[ 15.500558 0:10 fatfs::dir:140] Is a directory
[ 15.502442 0:10 fatfs::dir:140] Is a directory
[ 15.504179 0:10 fatfs::dir:140] Is a directory
[ 15.505942 0:10 fatfs::dir:140] Is a directory
[ 15.508128 0:10 fatfs::dir:140] Is a directory
[ 15.510001 0:10 fatfs::dir:140] Is a directory
[ 15.510001 0:10 fatfs::dir:140] Is a directory
[ 15.548939 0:10 fatfs::dir:140] Is a directory
1970-01-01 00:00:15.563 D [MediaServer] [9-MediaServer] System.cpp:133 startDaemon | 启动子进程:13
[ 15.589967 0:10 fatfs::dir:140] Is a directory
[ 15.591415 0:10 fatfs::dir:140] Is a directory
1970-01-01 00:00:15.610 I [MediaServer] [13-MediaServer] System.cpp:177 systemSetup | core文件大小设置为:18446744073709551615
[ 15.646279 0:14 fatfs::dir:140] Is a directory
[ 15.647834 0:14 fatfs::dir:140] Is a directory
1970-01-01 00:00:15.664 I [MediaServer] [13-MediaServer] System.cpp:186 systemSetup | 文件最大描述符个数设置为:18446744073709551615
QEMU: Terminated

```

添加之后执行结果为



```

11.332328 0:8 axstarry::syscall:38] [syscall] id = SIGACTION, args = [22, 1073739808, 0, 8, 11056, 1], entry
11.334617 0:8 axstarry::syscall_task::imp::signal:20] signum: 22, action: 3FFFF820, old action: 0
11.336530 0:8 axstarry::syscall:51] [syscall] id = 13, args = [22, 1073739808, 0, 8, 11056, 1], return 0
11.338578 0:8 axstarry::syscall:38] [syscall] id = SIGACTION, args = [2, 1073739808, 0, 8, 11056, 1], entry
11.341199 0:8 axstarry::syscall_task::imp::signal:20] signum: 2, action: 3FFFF820, old action: 0
11.343150 0:8 axstarry::syscall:51] [syscall] id = 13, args = [2, 1073739808, 0, 8, 11056, 1], return 0
11.345361 0:8 axstarry::syscall:38] [syscall] id = SIGACTION, args = [15, 1073739808, 0, 8, 11056, 1], entry
11.347451 0:8 axstarry::syscall_task::imp::signal:20] signum: 15, action: 3FFFF820, old action: 0
11.349383 0:8 axstarry::syscall:51] [syscall] id = 13, args = [15, 1073739808, 0, 8, 11056, 1], return 0
11.351638 0:8 axstarry::syscall:38] [syscall] id = SIGACTION, args = [3, 1073739808, 0, 8, 11056, 1], entry
11.354755 0:8 axstarry::syscall_task::imp::signal:20] signum: 3, action: 3FFFF820, old action: 0
11.356658 0:8 axstarry::syscall:51] [syscall] id = 13, args = [3, 1073739808, 0, 8, 11056, 1], return 0
11.360186 0:8 axstarry::syscall:38] [syscall] id = EXECVE, args = [6360, 6424, 6448, 11420, 3399988123389603631, 9259542123273814144], entry
13.786608 0:8 apxprocess::api:192] args: ["/lib64/ld-linux-x86-64.so.2", "/MediaServer", "-d"]
13.789337 0:8 apxprocess::api:194] The elf base addr may be different in different arch!
13.791129 0:8 elf_parser:142] Base addr for the elf: 0x4000000
13.792670 0:8 elf_parser:66] Base addr for the elf: 0x4000000
13.794158 0:8 elf_parser::arch::x86_64:55] Base addr for the elf: 0x4000000
13.795863 0:8 elf_parser::arch::x86_64:69] Relocating .rela.dyn
13.797633 0:8 elf_parser::arch::x86_64:154] Relocating .rela.plt
13.799085 0:8 elf_parser::arch::x86_64:178] Relocating done
13.819216 0:8 apxprocess::api:226] [new region] user heap: [VA:0x3fa00000, VA:0x3fe00000]
13.820990 0:8 elf_parser::auxv:48] ELF header addr: 0x4000000
13.833430 0:8 apxprocess::api:245] [new region] user stack: [VA:0x3fe00000, VA:0x40000000]
13.836721 0:8 axstarry::syscall:51] [syscall] id = 59, args = [6360, 6424, 6448, 11420, 3399988123389603631, 9259542123273814144], return 2
13.843824 0:8 axstarry::syscall:20] [syscall] id = BRK, args = [0, 1073740732, 0, 35, 0, 0], entry
13.847119 0:8 axstarry::syscall:51] [syscall] id = 12, args = [0, 1073740732, 0, 35, 0, 0], return 1067450368
13.850437 0:8 axstarry::syscall:38] [syscall] id = ARCH_PRCTL, args = [12289, 1073740656, 67236800, 1, 3, 2048], entry
13.853545 0:8 axstarry::syscall:51] [syscall] id = 158, args = [12289, 1073740656, 67236800, 1, 3, 2048], return -22
13.856930 0:8 axstarry::syscall:29] [syscall] id = OPENAT, args = [4294967196, 67351264, 524288, 0, 524288, 67351264], entry
13.860827 0:8 axstarry::syscall_fs::imp::io:407] path: "/MediaServer"
13.873057 0:8 apxprocess::link:243] create_link: /MediaServer -> /MediaServer
13.874599 0:8 axstarry::syscall:51] [syscall] id = 257, args = [4294967196, 67351264, 524288, 0, 524288, 67351264], return 4
13.877181 0:8 axstarry::syscall:29] [syscall] id = READ, args = [4, 1073739128, 832, 0, 524288, 67351264], entry
13.879606 0:8 axstarry::syscall_fs::imp::io:32] [read()] fd: 4, buf: 0x3ffff578, len: 832
13.882102 0:8 axstarry::syscall:51] [syscall] id = 0, args = [4, 1073739128, 832, 0, 524288, 67351264], return 832
13.884899 0:8 axstarry::syscall:29] [syscall] id = PREAD64, args = [4, 1073738112, 784, 64, 1073738112, 67351264], entry
13.908408 0:8 axstarry::syscall:51] [syscall] id = 17, args = [4, 1073738112, 784, 64, 1073738112, 67351264], return 784
13.911622 0:8 axstarry::syscall:29] [syscall] id = PREAD64, args = [4, 1073738048, 48, 880, 1073738112, 0], entry
13.923982 0:8 axstarry::syscall:51] [syscall] id = 17, args = [4, 1073738048, 48, 880, 1073738112, 0], return 48
13.926326 0:8 axstarry::syscall:29] [syscall] id = PREAD64, args = [4, 1073737968, 68, 928, 1073738112, 0], entry
13.938559 0:8 axstarry::syscall:51] [syscall] id = 17, args = [4, 1073737968, 68, 928, 1073738112, 0], return 68
13.941761 0:8 axstarry::syscall:38] [syscall] id = GETPID, args = [1073737284, 0, 67307088, 67307088, 67351264, 1], entry
13.943868 0:8 axstarry::syscall:51] [syscall] id = 39, args = [1073737284, 0, 67307088, 67307088, 67351264, 1], return 7
13.946927 0:8 axstarry::syscall:29] [syscall] id = WRITEV, args = [2, 1073737296, 6, 4294967295, 1073737209, 4294967295], entry
13.949184 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x3fffee44, len: 12
7: [ 13.951298 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x4030650, len: 5
files[ 13.953399 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x3ffffe21, len: 13
./MediaServer[ 13.955664 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x4030657, len: 2
[[ 13.957594 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x3fffedf8, len: 1
0[ 13.959539 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x403065c, len: 24
]; generating link map

```

不同1:

zlm

```

[ 16.687586 0:6 axstarry::syscall_task::imp::task:429] not support setpgid, try to set 7
[ 16.690262 0:6 axstarry::syscall:51] [syscall] id = 109, args = [7, 7, 0, 8, 11056, 1], return 0
[ 16.693719 0:6 axstarry::syscall:20] [syscall] id = MMAP, args = [0, 4096, 3, 34, 18446744073709551615, 0], entry
[ 16.697987 0:6 axmem:269] [mmap] vaddr: [VA:0x0, VA:0x1000], MappingFlags(READ | WRITE | USER), fixed: false, backend: false
[ 16.701751 0:6 axmem:287] find free area
[ 16.703342 0:6 axmem:292] found area [VA:0x6000, VA:0x7000]
[ 16.705622 0:6 axstarry::syscall:51] [syscall] id = 9, args = [0, 4096, 3, 34, 18446744073709551615, 0], return 24576
[ 16.708979 0:6 axhal::arch::x86_64::trap:21] User #PF @ 0x40c8f88, fault_vaddr=0x6008, error_code=0x4
[ 16.712026 0:6 axstarry::syscall:20] [syscall] id = MUNMAP, args = [24576, 4096, 4096, 1, 24576, 1], entry
[ 16.715053 0:6 axmem:310] [munmap] [VA:0x6000, VA:0x7000]
[ 16.716839 0:6 axmem:181] splitting for [VA:0x6000, VA:0x7000]
[ 16.718729 0:6 axmem:186] drop [VA:0x6000, VA:0x7000]
[ 16.720599 0:6 axstarry::syscall:51] [syscall] id = 11, args = [24576, 4096, 4096, 1, 24576, 1], return 0
[ 16.723889 0:6 axstarry::syscall:29] [syscall] id = IOCTL, args = [0, 21505, 1073739960, 9196, 0, 2], entry
[ 16.727007 0:6 axstarry::syscall_fs::imp::ctl:438] fd: 0, request: 21505, argp: 1073739960
[ 16.729913 0:6 apxprocess::stdio:123] stdin TCGETS | TIOCSPPGRP, pretend to be tty.
[ 16.733201 0:6 axstarry::syscall:51] [syscall] id = 16, args = [0, 21505, 1073739960, 9196, 0, 2], return 0
[ 16.736702 0:6 axstarry::syscall:20] [syscall] id = MMAP, args = [0, 16384, 3, 34, 18446744073709551615, 0], entry
[ 16.740214 0:6 axmem:269] [mmap] vaddr: [VA:0x0, VA:0x4000], MappingFlags(READ | WRITE | USER), fixed: false, backend: false
[ 16.744637 0:6 axmem:287] find free area
[ 16.746644 0:6 axmem:292] found area [VA:0x6000, VA:0xa000]
[ 16.748909 0:6 axstarry::syscall:51] [syscall] id = 9, args = [0, 16384, 3, 34, 18446744073709551615, 0], return 24576
[ 16.752626 0:6 axhal::arch::x86_64::trap:21] User #PF @ 0x40c8f88, fault_vaddr=0x6008, error_code=0x4
[ 16.756360 0:6 axhal::arch::x86_64::trap:21] User #PF @ 0x40c864d, fault_vaddr=0x7080, error_code=0x6

```

sleep

```

[ 18.190262 0:6 axstarry::syscall_task::imp::task:429] not support setpgid, try to set 7
[ 18.192819 0:6 axstarry::syscall:51] [syscall] id = 109, args = [7, 7, 0, 8, 11056, 1], return 0
[ 18.196160 0:6 axstarry::syscall:29] [syscall] id = IOCTL, args = [0, 21505, 1073739960, 9196, 0, 2], entry
[ 18.199199 0:6 axstarry::syscall_fs::imp::ctl:438] fd: 0, request: 21505, argp: 1073739960
[ 18.201869 0:6 apxprocess::stdio:123] stdin TCGETS | TIOCSPPGRP, pretend to be tty.
[ 18.204472 0:6 axstarry::syscall:51] [syscall] id = 16, args = [0, 21505, 1073739960, 9196, 0, 2], return 0
[ 18.207581 0:6 axstarry::syscall:20] [syscall] id = MMAP, args = [0, 16384, 3, 34, 18446744073709551615, 0], entry
[ 18.210931 0:6 axmem:269] [mmap] vaddr: [VA:0x0, VA:0x4000], MappingFlags(READ | WRITE | USER), fixed: false, backend: false
[ 18.214731 0:6 axmem:287] find free area
[ 18.216335 0:6 axmem:292] found area [VA:0x6000, VA:0xa000]
[ 18.218169 0:6 axstarry::syscall:51] [syscall] id = 9, args = [0, 16384, 3, 34, 18446744073709551615, 0], return 24576
[ 18.221080 0:6 axhal::arch::x86_64::trap:21] User #PF @ 0x40c8f88, fault_vaddr=0x6008, error_code=0x4
[ 18.223721 0:6 axhal::arch::x86_64::trap:21] User #PF @ 0x40c864d, fault_vaddr=0x7070, error_code=0x6

```

不同2:

zlm



```
[ 19.672786 0:8 axstarry::syscall:51] [syscall] id = 20, args = [2, 1073737296, 6, 4294967295, 1073737209, 4294967295], return 57
[ 19.677989 0:8 axstarry::syscall:29] [syscall] id = GETCWD, args = [67352736, 128, 18446744073642198880, 4294967295, 536870912, 67352696], entry
[ 19.681952 0:8 axstarry::syscall:51] [syscall] id = 79, args = [67352736, 128, 18446744073642198880, 4294967295, 536870912, 67352696], return 67352736
[ 19.686604 0:8 axstarry::syscall:29] [syscall] id = PREAD64, args = [4, 1073737808, 784, 64, 536870912, 67352696], entry
[ 19.700459 0:8 axstarry::syscall:51] [syscall] id = 17, args = [4, 1073737808, 784, 64, 536870912, 67352696], return 784
[ 19.704326 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [0, 20314104, 1, 2050, 4, 0], entry
[ 19.707530 0:8 axmem:269] [mmap] vaddr: [VA:0x0, VA:0x1360000], MappingFlags(READ | USER), fixed: false, backend: true
[ 19.711728 0:8 axmem:287] find free area
[ 19.713503 0:8 axmem:292] found area [VA:0x1000, VA:0x1361000]
[ 19.716008 0:8 axstarry::syscall:51] [syscall] id = 9, args = [0, 20314104, 1, 2050, 4, 0], return 4096
[ 19.719024 0:8 axstarry::syscall:20] [syscall] id = MPROTECT, args = [10264576, 9883648, 0, 2050, 4, 0], entry
[ 19.723078 0:8 axmem:340] [mprotect] addr: [VA:0x9ca000, VA:0x1337000], flags: MappingFlags(USER)
[ 19.727046 0:8 axstarry::syscall:51] [syscall] id = 10, args = [10264576, 9883648, 0, 2050, 4, 0], return 0
[ 19.730931 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [10264576, 6320128, 5, 2066, 4, 10260480], entry
[ 19.734884 0:8 axmem:269] [mmap] vaddr: [VA:0x9ca000, VA:0x1337000], MappingFlags(READ | EXECUTE | USER), fixed: true, backend: true
[ 19.738731 0:8 axmem:181] splitting for [VA:0x9ca000, VA:0x1337000]
[ 19.741116 0:8 axmem:204] shrink left [VA:0x9ca000, VA:0x1337000] to [VA:0x1337000, VA:0x1337000]
[ 19.744368 0:8 axstarry::syscall:51] [syscall] id = 9, args = [10264576, 6320128, 5, 2066, 4, 10260480], return 10264576
[ 19.748262 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [16584704, 3559424, 1, 2066, 4, 16580608], entry
[ 19.751684 0:8 axmem:269] [mmap] vaddr: [VA:0x1337000, VA:0x1336000], MappingFlags(READ | USER), fixed: true, backend: true
[ 19.755631 0:8 axmem:181] splitting for [VA:0x1337000, VA:0x1336000]
[ 19.758018 0:8 axmem:204] shrink left [VA:0x1337000, VA:0x1336000] to [VA:0x1336000, VA:0x1337000]
[ 19.761006 0:8 axstarry::syscall:51] [syscall] id = 9, args = [16584704, 3559424, 1, 2066, 4, 16580608], return 16584704
[ 19.764625 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [20148224, 139264, 3, 2066, 4, 20140032], entry
[ 19.767833 0:8 axmem:269] [mmap] vaddr: [VA:0x1337000, VA:0x1359000], MappingFlags(READ | WRITE | USER), fixed: true, backend: true
[ 19.771543 0:8 axmem:181] splitting for [VA:0x1337000, VA:0x1359000]
[ 19.773604 0:8 axmem:204] shrink left [VA:0x1337000, VA:0x1361000] to [VA:0x1359000, VA:0x1361000]
[ 19.776842 0:8 axstarry::syscall:51] [syscall] id = 9, args = [20148224, 139264, 3, 2066, 4, 20140032], return 20148224
[ 19.780091 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x402b1da, fault_vaddr=0x1358760, error_code=0x6
[ 19.796847 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [20287488, 30712, 3, 50, 4294967295, 0], entry
[ 19.800202 0:8 axmem:269] [mmap] vaddr: [VA:0x1359000, VA:0x1361000], MappingFlags(READ | WRITE | USER), fixed: true, backend: false
[ 19.804968 0:8 axmem:181] splitting for [VA:0x1359000, VA:0x1361000]
[ 19.807336 0:8 axmem:186] drop [VA:0x1359000, VA:0x1361000]
[ 19.810250 0:8 axstarry::syscall:51] [syscall] id = 9, args = [20287488, 30712, 3, 50, 4294967295, 0], return 20287488
[ 19.814219 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x4009562, fault_vaddr=0x1355f40, error_code=0x4
[ 19.830642 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x40095ad, fault_vaddr=0x1356000, error_code=0x4
[ 19.845999 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x400983b, fault_vaddr=0x1318, error_code=0x4
```

sleep

```
[ 18.569402 0:8 axstarry::syscall:51] [syscall] id = 20, args = [2, 1073737312, 6, 4294967295, 1073737225, 4294967295], return 51
[ 18.573629 0:8 axstarry::syscall:29] [syscall] id = GETCWD, args = [67352736, 128, 18446744073642198880, 4294967295, 536870912, 67352696], entry
[ 18.577891 0:8 axstarry::syscall:51] [syscall] id = 79, args = [67352736, 128, 18446744073642198880, 4294967295, 536870912, 67352696], return 67352736
[ 18.583039 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [0, 16408, 1, 2050, 4, 0], entry
[ 18.586022 0:8 axmem:269] [mmap] vaddr: [VA:0x0, VA:0x5000], MappingFlags(READ | USER), fixed: false, backend: true
[ 18.589720 0:8 axmem:287] find free area
[ 18.591251 0:8 axmem:292] found area [VA:0x1000, VA:0x6000]
[ 18.593403 0:8 axstarry::syscall:51] [syscall] id = 9, args = [0, 16408, 1, 2050, 4, 0], return 4096
[ 18.596713 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [8192, 4096, 5, 2066, 4, 4096], entry
[ 18.599782 0:8 axmem:269] [mmap] vaddr: [VA:0x2000, VA:0x3000], MappingFlags(READ | EXECUTE | USER), fixed: true, backend: true
[ 18.603905 0:8 axmem:181] splitting for [VA:0x2000, VA:0x3000]
[ 18.606337 0:8 axmem:191] split [VA:0x1000, VA:0x6000] into 2 areas
[ 18.609192 0:8 axstarry::syscall:51] [syscall] id = 9, args = [8192, 4096, 5, 2066, 4, 4096], return 8192
[ 18.612510 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [12288, 4096, 1, 2066, 4, 8192], entry
[ 18.615600 0:8 axmem:269] [mmap] vaddr: [VA:0x3000, VA:0x4000], MappingFlags(READ | USER), fixed: true, backend: true
[ 18.619575 0:8 axmem:181] splitting for [VA:0x3000, VA:0x4000]
[ 18.620959 0:8 axmem:204] shrink left [VA:0x3000, VA:0x6000] to [VA:0x4000, VA:0x6000]
[ 18.623701 0:8 axstarry::syscall:51] [syscall] id = 9, args = [12288, 4096, 1, 2066, 4, 8192], return 12288
[ 18.627041 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [16384, 8192, 3, 2066, 4, 8192], entry
[ 18.630245 0:8 axmem:269] [mmap] vaddr: [VA:0x4000, VA:0x6000], MappingFlags(READ | WRITE | USER), fixed: true, backend: true
[ 18.633844 0:8 axmem:181] splitting for [VA:0x4000, VA:0x6000]
[ 18.635906 0:8 axmem:186] drop [VA:0x4000, VA:0x6000]
[ 18.637992 0:8 axstarry::syscall:51] [syscall] id = 9, args = [16384, 8192, 3, 2066, 4, 8192], return 16384
[ 18.641433 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x402b260, fault_vaddr=0x5010, error_code=0x6
[ 18.645836 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x4009562, fault_vaddr=0x4dc0, error_code=0x4
[ 18.649392 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x400983b, fault_vaddr=0x12e0, error_code=0x4
```

不同3:

zlm

```
[ 19.943974 0:8 axstarry::syscall:51] [syscall] id = 20, args = [2, 1073736432, 18, 16, 1073736105, 4294967295], return 203
[ 19.948316 0:8 axstarry::syscall:38] [syscall] id = UNAME, args = [1073739600, 67111864, 0, 67271712, 0, 0], entry
[ 19.951814 0:8 axstarry::syscall:51] [syscall] id = 63, args = [1073739600, 67111864, 0, 67271712, 0, 0], return 0
[ 19.956858 0:8 axstarry::syscall:29] [syscall] id = ACCESS, args = [67304848, 4, 20275232, 0, 1, 67351280], entry
[ 19.960590 0:8 axstarry::syscall:51] [syscall] id = 21, args = [67304848, 4, 20275232, 0, 1, 67351280], return -2
[ 19.964699 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x402a837, fault_vaddr=0x1e8e6b, error_code=0x4
[ 19.980360 0:8 axstarry::syscall:38] [syscall] id = GETPID, args = [1073736020, 0, 67307488, 67307488, 0, 1], entry
[ 19.983742 0:8 axstarry::syscall:51] [syscall] id = 39, args = [1073736020, 0, 67307488, 67307488, 0, 1], return 7
[ 19.987257 0:8 axstarry::syscall:29] [syscall] id = WRITEV, args = [2, 1073736032, 12, 4294967295, 1073735897, 4294967295], entry
[ 19.991053 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x3fffe954, len: 12
7: [ 19.994110 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x40307e0, len: 1
```

sleep

```
[ 18.733193 0:8 axstarry::syscall:51] [syscall] id = 20, args = [2, 1073737248, 18, 16, 1073736921, 4294967295], return 203
[ 18.738188 0:8 axstarry::syscall:38] [syscall] id = UNAME, args = [1073739616, 67111864, 0, 18446744073709547520, 0, 0], entry
[ 18.742388 0:8 axstarry::syscall:51] [syscall] id = 63, args = [1073739616, 67111864, 0, 18446744073709547520, 0, 0], return 0
[ 18.747942 0:8 axstarry::syscall:29] [syscall] id = ACCESS, args = [67304848, 4, 20032, 0, 0, 67351280], entry
[ 18.751797 0:8 axstarry::syscall:51] [syscall] id = 21, args = [67304848, 4, 20032, 0, 0, 67351280], return -2
[ 18.755494 0:8 axstarry::syscall:38] [syscall] id = GETPID, args = [1073736036, 0, 67307488, 67307488, 0, 1], entry
[ 18.758693 0:8 axstarry::syscall:51] [syscall] id = 39, args = [1073736036, 0, 67307488, 67307488, 0, 1], return 7
[ 18.762065 0:8 axstarry::syscall:29] [syscall] id = WRITEV, args = [2, 1073736048, 12, 4294967295, 1073735913, 4294967295], entry
[ 18.766100 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x3fffe964, len: 12
7: [ 18.769332 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x40307e0, len: 1
```

不同4:

zlm

```
[ 20.157780 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x3fffe7c4, len: 12
7: [ 20.160883 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x4030650, len: 5
file=[ 20.164070 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x1e8e6b, len: 11
libssl.so.3[ 20.167547 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x4030657, len: 2
[[ 20.170446 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x3fffe778, len: 1
0[ 20.173156 0:8 axstarry::syscall_fs::imp::io:104] [write()] fd: 2, buf: 0x403065c, len: 24
]; generating link map
[ 20.176651 0:8 axstarry::syscall:51] [syscall] id = 20, args = [2, 1073735632, 6, 4294967295, 1073735545, 4294967295], return 55
[ 20.180416 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [0, 669680, 1, 2050, 4, 0], entry
[ 20.183420 0:8 axmem:269] [mmap] vaddr: [VA:0x0, VA:0xa4000], MappingFlags(READ | USER), fixed: false, backend: true
[ 20.186825 0:8 axmem:287] find free area
[ 20.188507 0:8 axmem:292] found area [VA:0x1361000, VA:0x1405000]
[ 20.190654 0:8 axstarry::syscall:51] [syscall] id = 9, args = [0, 669680, 1, 2050, 4, 0], return 20320256
[ 20.193890 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [20443136, 372736, 5, 2066, 4, 122880], entry
[ 20.197325 0:8 axmem:269] [mmap] vaddr: [VA:0x137f000, VA:0x13da000], MappingFlags(READ | EXECUTE | USER), fixed: true, backend: true
[ 20.201391 0:8 axmem:181] splitting for [VA:0x137f000, VA:0x13da000]
[ 20.203913 0:8 axmem:191] split [VA:0x1361000, VA:0x1405000] into 2 areas
[ 20.206681 0:8 axstarry::syscall:51] [syscall] id = 9, args = [20443136, 372736, 5, 2066, 4, 122880], return 20443136
[ 20.210242 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [20815872, 118784, 1, 2066, 4, 495616], entry
[ 20.213426 0:8 axmem:269] [mmap] vaddr: [VA:0x13da000, VA:0x13f7000], MappingFlags(READ | USER), fixed: true, backend: true
[ 20.217037 0:8 axmem:181] splitting for [VA:0x13da000, VA:0x13f7000]
[ 20.219026 0:8 axmem:204] shrink_left [VA:0x13da000, VA:0x1405000] to [VA:0x13f7000, VA:0x1405000]
[ 20.222157 0:8 axstarry::syscall:51] [syscall] id = 9, args = [20815872, 118784, 1, 2066, 4, 495616], return 20815872
[ 20.225927 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [20934656, 57344, 3, 2066, 4, 610304], entry
[ 20.229075 0:8 axmem:269] [mmap] vaddr: [VA:0x13f7000, VA:0x1405000], MappingFlags(READ | WRITE | USER), fixed: true, backend: true
[ 20.233172 0:8 axmem:181] splitting for [VA:0x13f7000, VA:0x1405000]
[ 20.235369 0:8 axmem:186] drop [VA:0x13f7000, VA:0x1405000]
[ 20.237351 0:8 axstarry::syscall:51] [syscall] id = 9, args = [20934656, 57344, 3, 2066, 4, 610304], return 20934656
[ 20.240831 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x402b1da, fault_vaddr=0x14047b0, error_code=0x6
[ 20.244530 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x4009562, fault_vaddr=0x13ffe30, error_code=0x4
[ 20.247765 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x40095ad, fault_vaddr=0x1400000, error_code=0x4
[ 20.251123 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x400983b, fault_vaddr=0x1361270, error_code=0x4
```

sleep

```
[ 19.166765 0:8 axstarry::syscall:51] [syscall] id = 20, args = [2, 1073734784, 18, 16, 1073734457, 4294967295], return 203
[ 19.170357 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x402723a, fault_vaddr=0x243ce, error_code=0x4
[ 19.174369 0:8 axstarry::syscall:20] [syscall] id = MMAP, args = [0, 8192, 3, 34, 4294967295, 0], entry
[ 19.178377 0:8 axmem:269] [mmap] vaddr: [VA:0x0, VA:0x2000], MappingFlags(READ | WRITE | USER), fixed: false, backend: false
[ 19.182008 0:8 axmem:287] find free area
[ 19.183603 0:8 axmem:292] found area [VA:0x22f000, VA:0x231000]
[ 19.185600 0:8 axstarry::syscall:51] [syscall] id = 9, args = [0, 8192, 3, 34, 4294967295, 0], return 2289664
[ 19.189081 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x4003a66, fault_vaddr=0x22f000, error_code=0x6
[ 19.192999 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x4015f72, fault_vaddr=0x25d08, error_code=0x4
[ 19.196605 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x4015faa, fault_vaddr=0x26010, error_code=0x4
[ 19.200408 0:8 axhal::arch::x86_64::trap:21] User #PF @ 0x4014462, fault_vaddr=0x2300c0, error_code=0x6
```

# MediaServer支持后台运行

答：Starry原本支持&后台运行，但是由于ZLM的问题导致Starry出现死循环问题，由输入判断futex存在问题，futex\_op为393，即

FUTEX\_WAIT\_BITSET | FUTEX\_PRIVATE\_FLAG | FUTEX\_CLOCK\_REALTIME

```
[syscall] id = FUTEX, args = [31251224, 393, 0, 0, 0, 4294967295], entry
[syscall] id = 202, args = [31251224, 393, 0, 0, 0, 4294967295], return -22
[syscall] id = FUTEX, args = [31251224, 393, 0, 0, 0, 4294967295], entry
[syscall] id = 202, args = [31251224, 393, 0, 0, 0, 4294967295], return -22
[syscall] id = FUTEX, args = [31251224, 393, 0, 0, 0, 4294967295], entry
[syscall] id = 202, args = [31251224, 393, 0, 0, 0, 0]FEMU: Terminated
```

解决方法：取消判断bitset和real\_futex\_val的按位与。

修改后仍然不能后台运行zlm，从日志中看到zlm持续占用

```

1970-01-01 00:00:41.695 I [MediaServer] [11-MediaServer] UdpServer.cpp:120 start_1 | UDP server bind to [::]: 10000
[ 42.861654 0:13 axstarry::syscall:51] [syscall] id = 1, args = [1, 31223488, 125, 1, 0, 134221232], return 125
[ 42.864904 0:13 axstarry::syscall:38] [syscall] id = GETPID, args = [31242376, 16708867, 27705056, 0, 32031296, 32156016], entry
[ 42.868664 0:13 axstarry::syscall:51] [syscall] id = 39, args = [31242376, 16708867, 27705056, 0, 32031296, 32156016], return 11
[ 42.872480 0:13 axstarry::syscall:29] [syscall] id = WRITE, args = [4, 31242336, 116, 1, 0, 32156416], entry
[ 42.875603 0:13 axstarry::syscall_fs::imp::io:104] [write()] fd: 4, buf: 0x1dcb860, len: 116
[ 42.878484 0:13 axstarry::syscall:51] [syscall] id = 1, args = [4, 31242336, 116, 1, 0, 32156416], return 116
[ 42.881872 0:13 axstarry::syscall:38] [syscall] id = GETPID, args = [31038064, 16708867, 27710320, 0, 32031296, 32156080], entry
[ 42.885824 0:13 axstarry::syscall:51] [syscall] id = 39, args = [31038064, 16708867, 27710320, 0, 32031296, 32156080], return 11
[ 42.889653 0:13 axstarry::syscall:29] [syscall] id = WRITE, args = [1, 31223488, 124, 1, 0, 134221280], entry
[ 42.893024 0:13 axstarry::syscall_fs::imp::io:104] [write()] fd: 1, buf: 0x1dc6ec0, len: 124
1970-01-01 00:00:42.198 I [MediaServer] [11-MediaServer] UdpServer.cpp:120 start_1 | UDP server bind to [::]: 9000
[ 42.899048 0:13 axstarry::syscall:51] [syscall] id = 1, args = [1, 31223488, 124, 1, 0, 134221280], return 124
[ 42.902794 0:13 axstarry::syscall:38] [syscall] id = GETPID, args = [31242376, 16708867, 27705056, 0, 32031296, 32156016], entry
[ 42.906959 0:13 axstarry::syscall:51] [syscall] id = 39, args = [31242376, 16708867, 27705056, 0, 32031296, 32156016], return 11
[ 42.910867 0:13 axstarry::syscall:29] [syscall] id = WRITE, args = [4, 31242336, 115, 1, 0, 32156416], entry
[ 42.913968 0:13 axstarry::syscall_fs::imp::io:104] [write()] fd: 4, buf: 0x1dcb860, len: 115
[ 42.916714 0:13 axstarry::syscall:51] [syscall] id = 1, args = [4, 31242336, 115, 1, 0, 32156416], return 115
[ 42.920037 0:13 axstarry::syscall:38] [syscall] id = FUTEX, args = [31251228, 393, 0, 0, 0, 4294967295], entry
[ 42.923721 0:13 axstarry::syscall_task::imp::futex:177] real val: 0x0, expected val: 0x0
[ 42.926715 0:15 axstarry::syscall:51] [syscall] id = 232, args = [7, 32677344, 1024, 4294967295, 0, 268438432], return 1
[ 42.930983 0:15 axstarry::syscall:29] [syscall] id = EPOLL_CTL, args = [7, 2, 0, 0, 0, 268438480], entry
[ 42.934231 0:15 axstarry::syscall:51] [syscall] id = 233, args = [7, 2, 0, 0, 0, 268438480], return -14
[ 42.937348 0:15 axstarry::syscall:29] [syscall] id = EPOLL_WAIT, args = [7, 32677344, 1024, 4294967295, 0, 268438528], entry
[ 42.940764 0:15 axstarry::syscall:51] [syscall] id = 232, args = [7, 32677344, 1024, 4294967295, 0, 268438528], return 1
[ 42.944363 0:15 axstarry::syscall:29] [syscall] id = EPOLL_CTL, args = [7, 2, 0, 0, 0, 268438576], entry
[ 42.947528 0:15 axstarry::syscall:51] [syscall] id = 233, args = [7, 2, 0, 0, 0, 268438576], return -14
[ 42.950643 0:15 axstarry::syscall:29] [syscall] id = EPOLL_WAIT, args = [7, 32677344, 1024, 4294967295, 0, 268438624], entry
[ 42.954285 0:15 axstarry::syscall:51] [syscall] id = 232, args = [7, 32677344, 1024, 4294967295, 0, 268438624], return 1
[ 42.957787 0:15 axstarry::syscall:29] [syscall] id = EPOLL_CTL, args = [7, 2, 0, 0, 0, 268438672], entry
[ 42.960984 0:15 axstarry::syscall:51] [syscall] id = 233, args = [7, 2, 0, 0, 0, 268438672], return -14
[ 42.964488 0:15 axstarry::syscall:29] [syscall] id = EPOLL_WAIT, args = [7, 32677344, 1024, 4294967295, 0, 268438720], entry
[ 42.968197 0:15 axstarry::syscall:51] [syscall] id = 232, args = [7, 32677344, 1024, 4294967295, 0, 268438720], return 1
[ 42.971765 0:15 axstarry::syscall:29] [syscall] id = EPOLL_CTL, args = [7, 2, 0, 0, 0, 268438768], entry
[ 42.974943 0:15 axstarry::syscall:51] [syscall] id = 233, args = [7, 2, 0, 0, 0, 268438768], return -14
[ 42.978006 0:15 axstarry::syscall:29] [syscall] id = EPOLL_WAIT, args = [7, 32677344, 1024, 4294967295, 0, 268438816], entry
[ 42.981690 0:15 axstarry::syscall:51] [syscall] id = 232, args = [7, 32677344, 1024, 4294967295, 0, 268438816], return 1
[ 42.985171 0:15 axstarry::syscall:29] [syscall] id = EPOLL_CTL, args = [7, 2, 0, 0, 0, 268438864], entry
[ 42.988224 0:15 axstarry::syscall:51] [syscall] id = 233, args = [7, 2, 0, 0, 0, 268438864], return -14
[ 42.991362 0:15 axstarry::syscall:29] [syscall] id = EPOLL_WAIT, args = [7, 32677344, 1024, 4294967295, 0, 268438912], entry
[ 42.994207 0:15 axstarry::syscall:51] [syscall] id = 232, args = [7, 32677344, 1024, 4294967295, 0, 268438912], return 1
[ 42.997767 0:15 axstarry::syscall:29] [syscall] id = EPOLL_CTL, args = [7, 2, 0, 0, 0, 268438960], entry
[ 43.000891 0:15 axstarry::syscall:51] [syscall] id = 233, args = [7, 2, 0, 0, 0, 268438960], return -14
[ 43.004271 0:15 axstarry::syscall:29] [syscall] id = EPOLL_WAIT, args = [7, 32677344, 1024, 4294967295, 0, 268439008], entry
[ 43.007844 0:15 axstarry::syscall:51] [syscall] id = 232, args = [7, 32677344, 1024, 4294967295, 0, 268439008], return 1
[ 43.011584 0:15 axstarry::syscall:29] [syscall] id = EPOLL_CTL, args = [7, 2, 0, 0, 0, 268439056], entry
[ 43.014817 0:15 axstarry::syscall:51] [syscall] id = 233, args = [7, 2, 0, 0, 0, 268439056], return -14
[ 43.018047 0:15 axstarry::syscall:29] [syscall] id = EPOLL_WAIT, args = [7, 32677344, 1024, 4294967295, 0, 268439104], entry
[ 43.021549 0:15 axstarry::syscall:51] [syscall] id = 232, args = [7, 32677344, 1024, 4294967295, 0, 268439104], return 1
[ 43.025055 0:15 axstarry::syscall:29] [syscall] id = EPOLL_CTL, args = [7, 2, 0, 0, 0, 268439152], entry

```

**原因：**这个版本的Starry没有支持抢占，导致15号任务拿着cpu不能释放

**解决方法：**Starry分离了一个新版本，支持多核和抢占，将代码修改移植到这个版本中，开启抢占特性

拉取<https://github.com/Starry-OS/Starry/>代码后在main分支下先执行

```

cargo update
cargo update --precise 0.4.19 log

```

然后测试这个版本是否可运行

```

make A=apps/monolithic_userboot LOG=off ACCEL=n FEATURES=img,sched_rr
APP_FEATURES=batch ARCH=x86_64 run

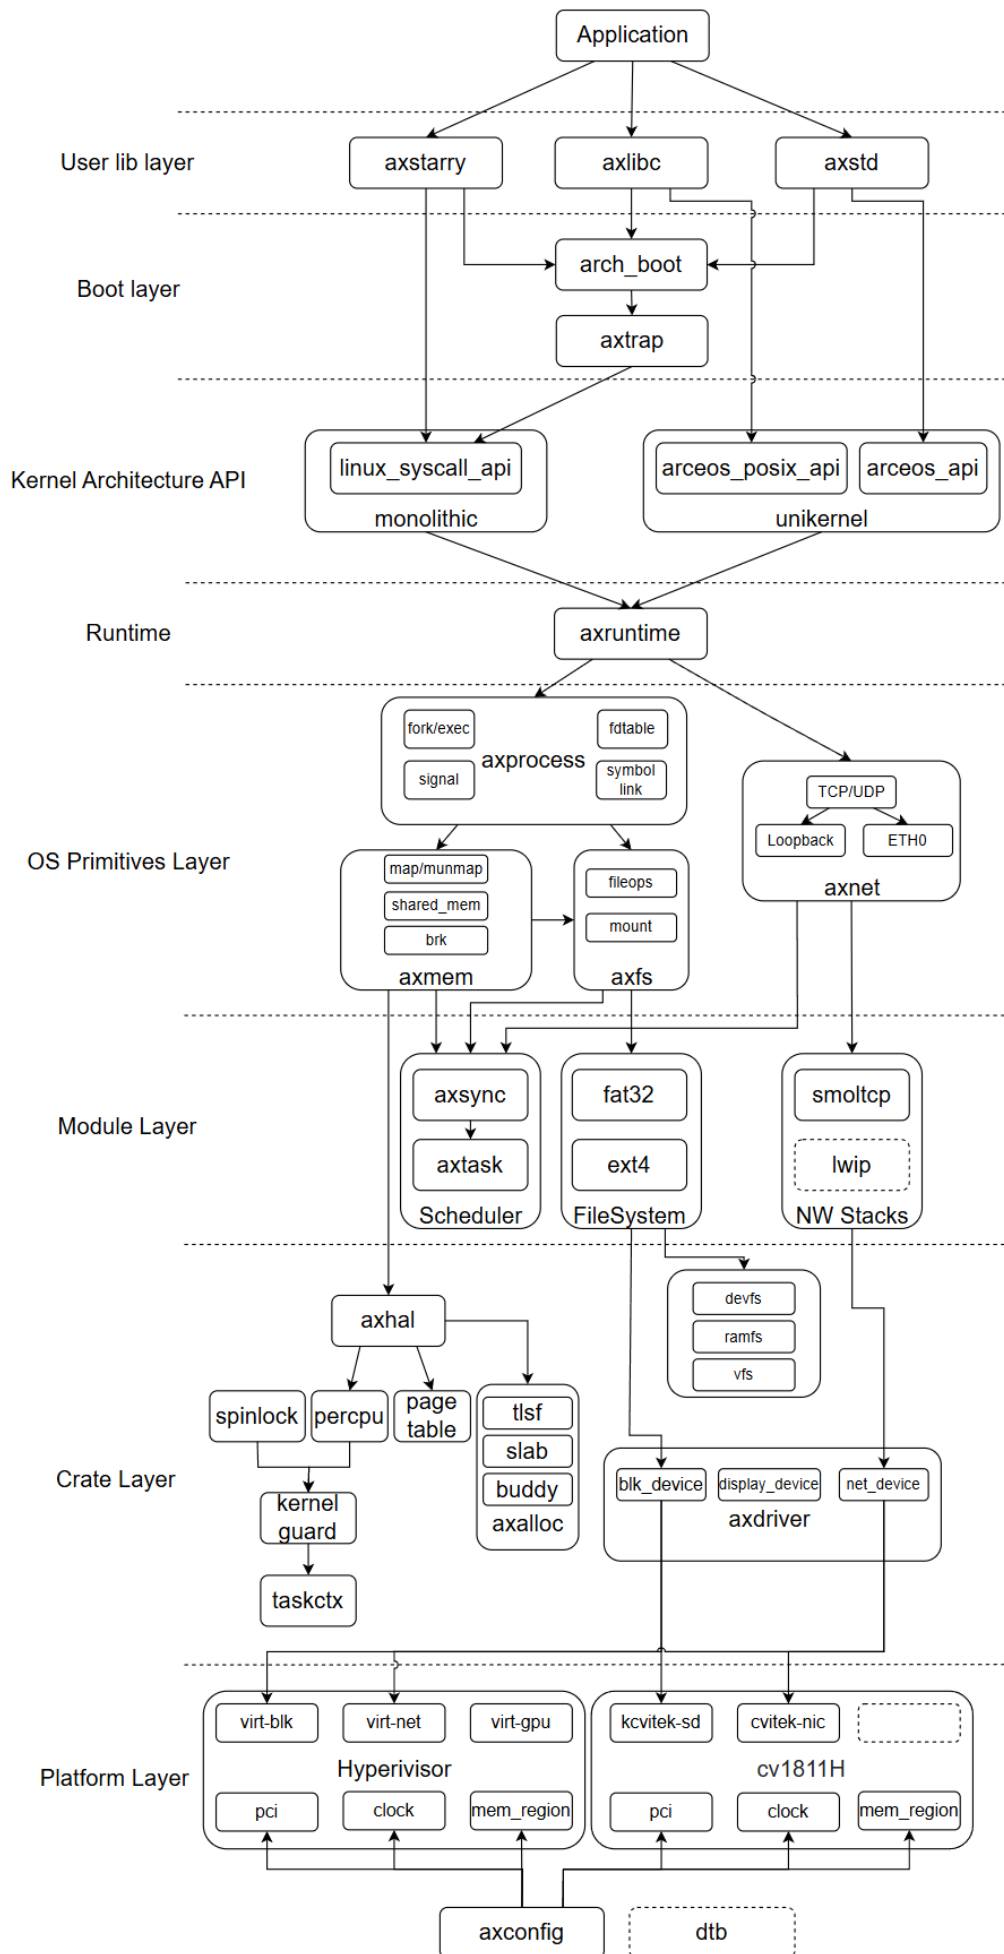
```

musl, libc 的测试集合没有问题，但是到最后的 libc-test 的部分测例会出现卡住的问题，不过不影响运行。

根据需要添加指定的模块



## Structure



## ZLM的话添加

```
kbuild patch add arceos_posix_api axruntime axfs axstarry linux_syscall_api
```

## 构建zlm的系统镜像

```
./build_img.sh -s 80 -file x86_64_ZLM -fs fat32
```

运行时加上multitask, sched\_rr特性, 其中img特性将virt-io切换成ramdisk, 提高内核运行速度

```
make LOG=error FEATURES=img,multitask,sched_rr,sched_cfs run
```

## 后台执行MediaServer后效果如下图所示

```
1970-01-01 00:00:21.119 I [MediaServer] [11-MediaServer] main.cpp:352 start_main | 已启动http hook 接口
1970-01-01 00:00:21.122 W [MediaServer] [11-MediaServer] main.cpp:360 start_main | The api.secret is invalid, modified it
1970-01-01 00:00:21.370 I [MediaServer] [11-MediaServer] TcpServer.cpp:218 start_l | TCP server listening on [::]: 554
1970-01-01 00:00:21.411 I [MediaServer] [11-MediaServer] TcpServer.cpp:218 start_l | TCP server listening on [::]: 332
1970-01-01 00:00:21.431 I [MediaServer] [11-MediaServer] TcpServer.cpp:218 start_l | TCP server listening on [::]: 1935
1970-01-01 00:00:21.462 I [MediaServer] [11-MediaServer] TcpServer.cpp:218 start_l | TCP server listening on [::]: 19350
1970-01-01 00:00:21.481 I [MediaServer] [11-MediaServer] TcpServer.cpp:218 start_l | TCP server listening on [::]: 80
1970-01-01 00:00:20.827 W [MediaServer] [11-stamp thread] util.cpp:385 operator() | Stamp expired is abnormal: 1085521
1970-01-01 00:00:21.641 I [MediaServer] [11-MediaServer] TcpServer.cpp:218 start_l | TCP server listening on [::]: 443
1970-01-01 00:00:21.662 I [MediaServer] [11-MediaServer] TcpServer.cpp:218 start_l | TCP server listening on [::]: 9000
1970-01-01 00:00:21.774 I [MediaServer] [11-MediaServer] TcpServer.cpp:218 start_l | TCP server listening on [::]: 10000
1970-01-01 00:00:21.826 I [MediaServer] [11-MediaServer] UdpServer.cpp:120 start_l | UDP server bind to [::]: 10000
1970-01-01 00:00:21.861 I [MediaServer] [11-MediaServer] UdpServer.cpp:120 start_l | UDP server bind to [::]: 9000

/, # ./busybox ps -a
PID USER TIME COMMAND
/ # ls
01.mp4 downloads.sh libcrypto.so.3 proc
MediaServer ffmpeg libgcc_s.so.1 sleep
busybox lat_sig libm.so.6 sleep.c
config.ini lib libssl.so.3 sys
crash.11 lib64 libstdc++.so.6 tmp
dev libc.so.6 log var
```